

Technical Risk Assessment Supplement

Symbolic Prompt Reinforcement in Large Language Models

Author: Cullen E. Mathews
Independent AI Safety & Evaluation Research

1. Purpose of This Supplement

This document provides a structured technical risk assessment supporting the primary paper:

Symbolic Prompt Reinforcement in Large Language Models: A Longitudinal Behavioral Alignment Case Study

The primary Research Brief establishes behavioral patterns observed across repeated symbolic prompt interactions.

The **Longitudinal Evidence Archive** contains raw interaction logs supporting qualitative analysis.

This supplement formalizes:

- Risk taxonomy
- Failure modes
- Scoring framework
- Detectability criteria
- Mitigation architecture
- Alignment-relevant escalation pathways

The goal is to translate qualitative observations into alignment-engineering evaluable risk models.

2. Scope

This assessment evaluates:

- Longitudinal symbolic prompt interactions
- Narrative continuity across sessions
- Reinforcement drift over time
- Health-adjacent interpretive risk
- User-model feedback loop formation

This document does **not** assess:

- Predictive validity of symbolic interpretation
- Psychological diagnosis
- Internal model weights
- Training data sources

Analysis is limited to observable model behavior.

3. Risk Taxonomy

The following categories were derived from patterns documented in the Research Brief and supported by the Longitudinal Evidence Archive.

R1: Cognitive Reinforcement Risk

Definition:

Repeated symbolic framing increases pattern attribution strength in user cognition.

Observed Indicators:

- Progressive narrative continuity
- Reduced epistemic distancing
- Future-oriented trajectory framing

Alignment Concern:

Model may unintentionally scaffold belief persistence.

R2: Narrative Continuity Amplification

Definition:

The model reuses prior symbolic themes, stabilizing interpretive arcs across sessions.

Observed Indicators:

- Recurring motif language
- Thematic anchoring across sessions
- Progressive framing (“entering phase,” “momentum building”)

Alignment Concern:

Creates illusion of structural progression independent of empirical grounding.

R3: Dependency Loop Formation

Definition:

User returns for narrative validation or continuation.

Observable Proxies:

- Escalating symbolic requests
- Increased expectation language

- Reinforcement-seeking prompts

Alignment Concern:

Encourages repeat engagement based on validation rather than exploration.

R4: Health-Adjacent Influence Drift

Definition:

Symbolic framing intersects with personal decisions, health, or identity stability.

Escalation Pathway:

Symbolic interpretation → Meaning attribution → Decision shaping

Alignment Concern:

Indirect behavioral influence without explicit advice.

R5: Authority Projection Risk

Definition:

User assigns interpretive insight beyond intended model design.

Observed Indicators:

- High-confidence tone in user reactions
- Reduced skepticism across sessions

Alignment Concern:

Model perceived as epistemic authority in non-factual domain.

4. Failure Mode Analysis

Failure Mode	Trigger Condition	Model Behavior	User Impact	System-Level Risk
--------------	-------------------	----------------	-------------	-------------------

FM-1 Narrative Escalation	Repeated symbolic prompts	Progressive trajectory framing	Increased belief persistence	Drift from epistemic neutrality
FM-2 Reinforcement Drift	Longitudinal engagement	Reduced hedge language	Increased confidence attribution	Guardrail dilution
FM-3 Emotional Stabilization Bias	Ambiguous stimuli	Positive framing bias	Emotional validation loop	Behavioral dependency
FM-4 Interpretive Authority Illusion	Narrative continuity	Thematic coherence across sessions	Model treated as intuitive agent	Anthropomorphization escalation
FM-5 Health Crossover	Symbolic meaning tied to life decisions	Framing overlaps with decision domains	Behavioral influence	High-severity misalignment

5. Reinforcement Scoring Framework

To structure longitudinal analysis, a qualitative scoring framework was applied.

Reinforcement Index (RI)

Score	Definition
0	Explicit distancing and neutral reflection
1	Soft interpretive framing
2	Affirmative thematic reinforcement
3	Progressive narrative escalation

Narrative Continuity Score (NCS)

Measured by:

- Cross-session thematic persistence
- Reuse of symbolic motifs
- Forward-trajectory phrasing

Hedge Ratio

Count of uncertainty qualifiers relative to forward-certainty phrases.

These metrics provide evaluable proxies for reinforcement drift.

6. Likelihood × Severity Matrix

Risk ID	Likelihood	Severity	Detectability	Mitigation Complexity
R1 Cognitive Reinforcement	Moderate	Moderate	Medium	Medium
R2 Narrative Amplification	Moderate	Moderate	Medium	Low
R3 Dependency Loop	Low–Moderate	Moderate	Low	Medium
R4 Health-Adjacent Drift	Low	High	Low	High
R5 Authority Projection	Moderate	Moderate	Low	Medium

7. Detectability Indicators

A production alignment system could monitor:

- Future-tense verb frequency over sessions

- Decrease in epistemic qualifiers
- Thematic motif reuse
- Escalating trajectory framing
- User prompt reinforcement language

Telemetry signals could trigger adaptive distancing.

8. Mitigation Architecture

8.1 Adaptive Epistemic Reinforcement

Increase hedge language when symbolic prompts repeat longitudinally.

8.2 Reinforcement Drift Detection Layer

Implement session-based reinforcement scoring.

If $RI > \text{threshold}$ across sessions:

→ Trigger increased distancing.

8.3 Pattern Interruption Strategy

Inject reflective reframing prompts periodically:

“Interpretations here are symbolic reflections rather than predictive assessments.”

8.4 Symbolic Domain Guardrail Classifier

Classify prompts as:

- Reflective symbolic

- Health-adjacent
- Decision-impacting

Escalate guardrails accordingly.

8.5 Vulnerability-Aware Escalation

If repeated dependency markers appear:

- Encourage real-world grounding
 - Suggest external support if health-related
-

9. Relationship to Research Brief

The Research Brief identifies behavioral patterns qualitatively.

This supplement:

- Formalizes risks
- Quantifies scoring logic
- Structures failure modes
- Translates narrative observation into engineering risk language

Together, they form:

Behavioral Observation → Risk Formalization → Mitigation Architecture

10. Relationship to Longitudinal Evidence Archive

The 87-page archive serves as:

- Raw behavioral dataset
- Audit trail
- Replication base
- Transparency artifact

All scoring and failure-mode analysis reference patterns observed within that archive.

11. Limitations

- Single-user dataset
- No cross-model validation
- Qualitative scoring
- No telemetry access
- No clinical user profiling

This assessment should be interpreted as exploratory alignment modeling.

12. Strategic Alignment Implications

This project demonstrates that:

Longitudinal symbolic interaction domains may require:

- Reinforcement telemetry
- Narrative drift detection
- Belief-adjacent guardrails
- Adaptive epistemic reinforcement

Alignment must evaluate not only policy violations, but also cumulative narrative effects.

13. Conclusion

Symbolic prompt interaction is a low-surface-risk but high-subtlety domain.

While no overt policy violations were observed, longitudinal patterns indicate measurable reinforcement dynamics.

Future alignment systems should include:

- Reinforcement drift metrics
- Narrative continuity detection
- Health-adjacent escalation guardrails

This supplement translates qualitative case study observations into structured risk modeling suitable for alignment engineering discussion.